FULL SEO AUDIT REPORT

japanesewithavnish.com

Prepared by Claude SEO

58/100

SEO Health Score

August 08, 2026 Full Audit

Table of Contents

1. Executive Summary

This report presents a comprehensive SEO audit of **japanesewithavnish.com** for a **EdTech / SaaS — subscription Japanese-language learning platform (JLPT N5–N1)** site, generated on August 08, 2026. Findings combine technical, content, schema, performance, visual, and search-readiness evidence as available.

58/100

SEO HEALTH SCORE

Critical Issues Found:

1. **Info:** CRITICAL: Every canonical, og:url, JSON-LD url and sitemap entry points to the default Netlify subdomain cosmic-blini-bc94d7.netlify.app instead of japanesewithavnish.com — this tells Google the netlify.app copy is the 'real' page and can de-index/dilute the primary domain.
2. **Info:** CRITICAL: The netlify.app subdomain serves the full site with HTTP 200 (live duplicate) instead of 301-redirecting to the primary domain.
3. **Info:** HIGH: robots.txt Sitemap directive and all 6,733 sitemap URLs use the wrong netlify.app domain.
4. **Info:** HIGH: Thousands of /learn/* and /blog/study_guide/* pages have auto-generated stub meta descriptions (just the word/title, e.g. '尻 - butt').
5. **Info:** MEDIUM: Homepage and most pages ship no og:image/twitter:image (broken social previews) despite twitter:card=summary_large_image.

Quick Wins:

1. Set Netlify env var NEXT_PUBLIC_SITE_URL=https://japanesewithavnish.com and redeploy — fixes canonical, og:url, sitemap, robots and JSON-LD in one change.
2. Add a 301 redirect from cosmic-blini-bc94d7.netlify.app to https://japanesewithavnish.com.
3. Populate JSON-LD Organization sameAs with the X/YouTube/Instagram profiles already linked in the footer, and add a logo.
4. Re-specify openGraph.images in page-level metadata so /logo.png (or a proper 1200x630 OG image) is not dropped.
5. Add a Google PageSpeed Insights API key so real Core Web Vitals can be measured.

2. Audit Categories

2.1 Technical SEO

Score: 45/100

What Works

- HTTP→HTTPS 301 and www→non-www 301 canonicalization both correct
- Proper 404 status on unknown URLs
- robots.txt present with sensible disallows (admin, api, account, checkout, login) and explicit rules for AI crawlers (OAI-SearchBot, ChatGPT-User)
- XML sitemap present and complete (6,733 URLs)
- Server-side rendered (not a client-only SPA) — content is in the initial HTML
- Security headers present: HSTS, X-Frame-Options, X-Content-Type-Options, Referrer-Policy, Permissions-Policy
- Netlify CDN with Brotli compression

Findings

Canonical tags site-wide point to the wrong (netlify.app) domain Critical

Every page returns <link rel="canonical" href="https://cosmic-blini-bc94d7.netlify.app/...">. Verified on /, /pricing, /start-here, /quiz, /jlpt, /blog and /learn/grammar. A canonical pointing to a different live domain instructs Google to index that domain instead of japanesewithavnish.com, which can drop the primary domain from search results and split any accumulated authority.

Recommendation: Root cause: dynamic metadata reads process.env.NEXT_PUBLIC_SITE_URL (src/lib/seo.ts:1, src/app/layout.tsx:47, src/app/sitemap.ts:4, src/app/robots.ts:3) which is set to the netlify.app URL in Netlify. Set NEXT_PUBLIC_SITE_URL=https://japanesewithavnish.com in Netlify env vars and redeploy.

Netlify subdomain serves a live duplicate of the whole site (HTTP 200) Critical

https://cosmic-blini-bc94d7.netlify.app/ returns 200 and renders the full site rather than redirecting. Combined with the canonical issue, Google sees two complete copies of every page and is being told the netlify.app copy is authoritative.

Recommendation: Add a 301 redirect from the *.netlify.app host to https://japanesewithavnish.com (Netlify redirect rule or middleware host check). This also protects against the subdomain being indexed independently.

Sitemap and robots.txt reference the wrong domain High

All 6,733 <loc> entries in sitemap.xml use cosmic-blini-bc94d7.netlify.app, and robots.txt declares 'Sitemap: https://cosmic-blini-bc94d7.netlify.app/sitemap.xml'. Search engines may treat these as cross-domain URLs and discount them.

Recommendation: Fixed automatically by the NEXT_PUBLIC_SITE_URL change (both are generated from the same env var in src/app/sitemap.ts and src/app/robots.ts). Re-submit the corrected sitemap in Google Search Console afterward.

HTML responses are marked non-cacheable Medium

Homepage Cache-Control is 'private,no-cache,no-store,max-age=0,must-revalidate'. Public marketing/lesson HTML should be cacheable at the edge with revalidation to improve TTFB and repeat-visit performance.

Recommendation: For public, non-personalized routes, serve a cacheable policy (e.g. 's-maxage' with stale-while-revalidate) instead of no-store. Keep no-store only for authenticated/account routes.

No Content-Security-Policy header Low

HSTS and other security headers are present, but there is no CSP. Not an SEO ranking factor, but relevant given third-party ad/analytics scripts (AdSense, Monetag, GTM) are loaded.

Recommendation: Add a report-only CSP first, then enforce once tuned for the known third parties.

2.2 On-Page SEO

Score: 52/100

What Works

- Homepage title is well-formed and unique: 'Japanese with Avnish | JLPT N5–N1 Learning' (43 chars)
- Titles are unique across the sampled pages (no duplicate-title clusters found)
- Exactly one H1 per page across the entire sample
- Clean, human-readable URL structure (/learn/grammar, /jlpt, /blog/...)
- Marketing-page meta descriptions are compelling and keyword-relevant

Findings

og:url and canonical use the wrong domain (see Technical) Critical

og:url on the homepage is https://cosmic-blini-bc94d7.netlify.app; shares and canonical signals resolve to the netlify.app domain.

Recommendation: Resolved by the NEXT_PUBLIC_SITE_URL fix.

Auto-generated stub meta descriptions at scale High

Thousands of programmatic pages use the word/title as the entire description: /learn/kanji/... => '尻 - butt', /learn/vocabulary/... => 'ときどき', and many /blog/study_guide/* pages have 24–65 character descriptions. These provide no differentiation and waste SERP snippet real estate.

Recommendation: Generate templated but descriptive meta descriptions per entity type, e.g. 'Learn the JLPT N2 kanji 尻 ("butt"): readings, stroke order, example words and practice.' Use the data you already have (level, readings, meaning).

Missing og:image / twitter:image on most pages Medium

The root layout sets openGraph.images=['/logo.png'], but the rendered homepage and ~33/39 sampled pages ship no og:image or twitter:image. Cause: page-level metadata objects redefine openGraph without re-specifying images, so Next.js drops the inherited image. twitter:card is 'summary_large_image' but there is no image to show.

Recommendation: Provide a proper 1200x630 default OG image and re-specify openGraph.images/twitter.images in every page-level metadata object (or centralize via a buildMetadata() helper in src/lib/seo.ts).

A few guide page titles exceed ~60 characters Low

e.g. 'Listening, Reading & Writing Practice — Site Guide | Japanese with Avnish' (77 chars) will be truncated in SERPs. The '— Site Guide | Japanese with Avnish' suffix consumes most of the budget.

Recommendation: Shorten the boilerplate suffix on /guide/* pages (e.g. drop 'Site Guide') so the descriptive part survives truncation.

2.3 Content Quality

Score: 68/100

What Works

- Lesson pages contain real, unique content — vocab pages include multiple example sentences with romaji + English (~380 words); kanji pages include readings, stroke count, frequency, study tips (~550 words)
- Study-guide posts are substantial and genuinely useful (~900 words, structured with clear sections)
- Blog demonstrates first-hand E-E-A-T (personal posts like 'my-journey-from-n5-to-n3', 'mistakes-i-made-while-learning-japanese')
- Strong, focused topical authority around JLPT N5–N1
- Content is fully crawlable (not gated behind login for the lesson detail pages)

Findings

Thin key conversion/landing pages **Medium**

Several important pages are light on indexable text: /quiz (143 words), /free-n5-pack (221 words), /learn hub (294 words), and several /guide/* pages (~260 words). These are entry points that should rank and convert.

Recommendation: Expand /quiz and /free-n5-pack with supporting copy (what the quiz measures, sample questions, what the free pack includes, FAQs). Add an intro + linked overview to the /learn hub.

Programmatic index-bloat risk on ~6,200 lesson pages **Medium**

5,230 vocabulary + 561 grammar + 424 kanji pages share a near-identical template. Individually defensible (unique word + examples), but at this scale with stub meta descriptions and templated boilerplate they risk being seen as thin/doorway pages.

Recommendation: Ensure each page has genuinely unique above-the-fold content, strengthen internal linking between related entries (same level, same topic), and consider consolidating the very thinnest single-item pages into level/topic hub pages with anchors.

Weak explicit E-E-A-T / author signals **Medium**

Organization JSON-LD has an empty sameAs and no logo; there is no Person/author schema or visible author credentials tying content to a named expert, despite the brand being built around 'Avnish'.

Recommendation: Add a Person/author entity (with credentials and sameAs to social profiles), author bylines on blog/lessons, and an About page that establishes expertise.

2.4 Schema / Structured Data

Score: **50/100**

What Works

- Valid Organization JSON-LD is present and parses correctly
- JSON-LD is delivered in the initial HTML (crawlable)

Findings

Organization schema is minimal and uses the wrong URL **Medium**

The only structured data is a single Organization block with url=netlify.app, sameAs=[] and no logo. It under-describes the entity.

Recommendation: Fix url via the env var change, add logo, and populate sameAs with the X/YouTube/Instagram profiles already linked in the footer.

High-value schema types are missing High

No Course/EducationalOrganization, BreadcrumbList, Article (blog), FAQPage (guide/jlpt pages), Product/Offer (pricing bundles), or WebSite+SearchAction. These are strong rich-result opportunities for an education/commerce site.

Recommendation: Add BreadcrumbList site-wide, Article to blog posts, Course/EducationalOccupationalProgram to JLPT level pages, Product+Offer to pricing bundles, FAQPage to guide/JLPT pages, and WebSite+SearchAction on the homepage.

2.5 Performance (Core Web Vitals)

Score: 60/100

What Works

- Served from Netlify CDN with Brotli compression
- Resource preload hints present (fonts)
- Server-side rendered content (fast first contentful paint of meaningful markup)
- Mobile layout renders cleanly and responsively

Findings

Field Core Web Vitals could not be measured Info

PageSpeed Insights / CrUX returned a rate-limit error (no Google API key configured), and the domain likely lacks a CrUX field dataset due to low traffic. Scores here are heuristic, not measured.

Recommendation: Add a PageSpeed Insights API key (claude-seo google_auth) and connect Google Search Console to obtain real LCP/INP/CLS field data.

Ad + analytics scripts add INP/CLS risk on a paid product Medium

The homepage loads Google AdSense (pagead2.googlesyndication) and Monetag alongside Google Tag Manager. Ad scripts are a common source of layout shift (CLS) and input-delay (INP), and ads on a premium subscription learning site can also undercut trust/conversion.

Recommendation: Reconsider running display-ad networks on a paid-subscription product. If kept, reserve fixed slot dimensions to prevent CLS, lazy-load below-the-fold ad units, and defer non-critical third-party JS.

Heavy hydration payload Low

Homepage HTML is ~131 KB including 16 inline Next.js hydration chunks and 16 script tags. Not alarming, but worth watching for INP on interaction-heavy lesson pages.

Recommendation: Audit client-component boundaries; keep marketing pages as light server components and measure INP on lesson pages once field data is available.

2.6 AI Search Readiness (GEO)

Score: 78/100

What Works

- Excellent llms.txt at /llms.txt — structured, descriptive, and (being a static file) correctly uses the japanesewithavnish.com domain
- robots.txt explicitly grants access to OAI-SearchBot and ChatGPT-User
- Server-side rendered content is fully readable by AI crawlers
- Clear, well-structured, question-oriented guide content that is easy to cite

Findings

Domain confusion may fracture AI citations Medium

llms.txt points to japanesewithavnish.com while canonical/og/schema point to netlify.app. AI engines that reconcile canonical signals may cite or attribute the wrong domain.

Recommendation: Resolved by the NEXT_PUBLIC_SITE_URL fix; keep llms.txt in sync with the primary domain.

Weak entity/authority signals for AI attribution Low

Empty sameAs and no Organization logo weaken the knowledge-graph entity that AI engines use to attribute and trust the brand.

Recommendation: Complete Organization + Person schema (logo, sameAs, founder) to strengthen entity recognition.

2.7 Images

Score: 62/100

What Works

- Images served from a CDN (Cloudflare R2)
- Favicon present
- Above-the-fold imagery renders crisply on desktop and mobile

Findings

~20% of sampled images lack alt text Medium

26 of 128 images across the sample had missing or empty alt attributes, hurting accessibility and image-search eligibility.

Recommendation: Add descriptive alt text to content images; keep purely decorative images with empty alt but proper role.

Missing dedicated social-share (OG) image Medium

No 1200x630 OG image exists; pages either drop the image entirely or fall back to a logo. Social/AI previews look bare.

Recommendation: Create a branded 1200x630 default OG image plus per-section variants (JLPT level, blog) and wire them into metadata.

Images served from a raw R2 dev hostname Low

Assets load from pub-8b7ba9883f454bddb9a3f2140d4aeab2.r2.dev rather than a branded custom domain.

Recommendation: Map a custom domain (e.g. cdn.japanesewithavnish.com) to the R2 bucket for branding and long-term URL stability.

3. Action Plan

Phase 1: Critical Fixes Week 1

- Set Netlify env var NEXT_PUBLIC_SITE_URL=https://japanesewithavnish.com and redeploy (fixes canonical, og:url, sitemap, robots, JSON-LD in one shot).
- Add a 301 redirect from cosmic-blini-bc94d7.netlify.app to https://japanesewithavnish.com.
- Verify canonicals on /, /pricing, /jlpt, a /learn/* and a /blog/* page now show japanesewithavnish.com.
- Re-submit the corrected sitemap in Google Search Console and request indexing of key pages.

Phase 2: High-Impact Improvements Weeks 2-3

- Generate descriptive, templated meta descriptions for /learn/* and /blog/study_guide/* pages using existing data (level, meaning, readings).
- Create a 1200x630 OG image and re-specify openGraph.images/twitter.images across page-level metadata (centralize in src/lib/seo.ts).
- Add BreadcrumbList site-wide, Article to blog posts, and Product+Offer to pricing bundles.
- Complete Organization schema (logo + sameAs) and add Person/author schema.
- Add descriptive alt text to content images.

Phase 3: Content & Authority Month 2

- Expand thin conversion pages: /quiz, /free-n5-pack, /learn hub, and /guide/* pages.
- Add Course/EducationalOccupationalProgram schema and FAQPage to JLPT level and guide pages.
- Strengthen internal linking between related lesson entries (same level/topic) to reduce thin-page risk.
- Reconsider or harden AdSense/Monetag on paid pages; reserve ad slot dimensions to protect CLS.
- Build out author/About pages and bylines to reinforce E-E-A-T.

Phase 4: Monitoring & Iteration Ongoing

- Add a PageSpeed Insights API key and connect Google Search Console + GA4 for real CWV, indexation and traffic data.
- Monitor indexation of japanesewithavnish.com vs the netlify.app subdomain in GSC until the duplicate is fully de-indexed.
- Track INP/CLS on lesson pages once field data accrues.
- Re-run this audit after Phase 1 deploys to confirm the canonical fix propagated.

Data Sources & Methodology

Source	Description	Update Frequency
PageSpeed Insights API	Lighthouse lab audit (mobile emulation, Moto G Power, slow 4G)	Real-time
Chrome UX Report (CrUX)	28-day rolling field data from real Chrome users	Daily ~04:00 UTC
CrUX History API	25-week p75 trend data per metric	Weekly
Google Search Console	Search Analytics (clicks, impressions, CTR, position)	2-3 day lag
URL Inspection API	Per-URL index status, coverage state, crawl info	Real-time (2,000/day)

Report generated by Claude SEO — Google SEO Intelligence Skill — August 08, 2026
Methodology based on Google Web Vitals thresholds, Search Console documentation, and Lighthouse scoring algorithms.